

Homework - 2: N-Queens

Adapted from the [Course Scheduling of Stanford CS221](#)

Introduction

In this project, you are going to construct a CSP for N-Queens problem. You will be given several python files.

Files you'll edit:

[submission.py](#) Where you need to fill in codes in three places.

Files you need to look at:

[grader.py](#) You can use this file to grade your program.

[util.py](#) Some important supporting functions are provided including the class CSP.

Files to Edit and Submit: You will fill in portions of [submission.py](#) during the assignment. You should submit this file with your code and comments. Please *do not* change the other files in this distribution or submit any of our original files other than these files.

Evaluation: Your code will be autograded for technical correctness. Please *do not* change the names of any provided functions or classes within the code, or you will wreak havoc on the autograder. However, the correctness of your implementation -- not the autograder's judgements -- will be the final judge of your score. If necessary, we will review and grade assignments individually to ensure that you receive due credit for your work.

Discussions: You can ask any questions about this homework through course email box. However, make sure you have already discussed it with your group members and cannot find any solution. **All inquiry emails related to this project should use a title starting with "hw-2".**

Academic Dishonesty: We will be checking your code against other submissions in the class for logical redundancy. If you copy someone else's code and submit it with minor changes, we will know. These cheat detectors are quite hard to fool, so please don't try. We trust you all to submit your own work only; *please* don't let us down. If you do, we will pursue the strongest consequences available to us.

CSP for N-queens

Notice we are already able to solve the CSPs, because in `submission.py`, a basic backtracking search is already implemented. The function will help you solve the CSP. You should read `BacktrackingSearch` carefully to make sure that you understand how the backtracking search is working.

Take a look at `BacktrackingSearch.reset_results()` to see the other fields which are set as a result of solving the CSP.

- a. [5 points] Let's create a CSP to solve the n-queens problem: Given an $n \times n$ board, we'd like to place n queens on this board such that no two queens are on the same row, column, or diagonal. Implement `create_nqueens_csp()` by adding n variables and some number of binary factors. Note that the solver collects some basic statistics on the performance of the algorithm. You should take advantage of these statistics for debugging and analysis. You should get 92 (optimal) assignments for $n=8$ with exactly 2057 operations (number of calls to `backtrack()`).

Hint: If you get a larger number of operations, make sure your CSP is minimal. Try to define the variables such that the size of domain is $O(n)$.

- b. [5 points] You might notice that our search algorithm explores quite a large number of states even for the 8×8 board. Let's see if we can do better. One heuristic we discussed in class is using most constrained variable (MCV): To choose an unassigned variable, pick the X_j that has the fewest number of values a which are consistent with the current partial assignment (a for which `get_delta_weight()` on $X_j=a$ returns a non-zero value). Implement this heuristic in `get_unassigned_variable()` under the condition `self.mcv = True`. It should take you exactly 1361 operations to find all optimal assignments for 8 queens CSP — that's 30% fewer!

Some useful fields:

- `csp.unaryFactors[var][val]` gives the unary factor value.
 - `csp.binaryFactors[var1][var2][val1][val2]` gives the binary factor value. Here, `var1` and `var2` are variables and `val1` and `val2` are their corresponding values.
 - In `BacktrackingSearch`, if `var` has been assigned a value, you can retrieve it using `assignment[var]`. Otherwise `var` is not in `assignment`.
- c. [10 points] The previous heuristics looked only at the local effects of a variable or value. Let's now implement arc consistency (AC-3) that we discussed in lecture. After we set variable X_j to value a , we remove the values b of all neighboring variables X_k that could cause arc-inconsistencies. If X_k 's domain has changed, we use X_k 's domain to remove values from the domains of its neighboring variables. This is repeated until no domains have changed. Note that this may significantly reduce your branching factor, although at some cost. In `backtrack()` we've implemented code which copies and restores domains for you. Your job is to fill in `arc_consistency_check()`.

With AC-3 enabled, it should take you 769 operations only to find all optimal assignments to 8 queens CSP — That is almost 45% fewer even compared with MCV!

Hint 1: documentation for `CSP.add_unary_factor()` and `CSP.add_binary_factor()` can be helpful.

Hint 2: although AC-3 works recursively, you may implement it iteratively. Using a queue might be a good idea. `li.pop(0)` removes and returns the first element for a python list `li`.

Submission (IMPORTANT THINGS! READ THREE TIMES 😊)

- Submit your code, **only one: submission.py**, through Course Emailbox in a zip file with the name **"hw2-n-queen"**.
- This is a homework so that you need to submit your own version of codes.
- You are encouraged to discuss with your group members since it might take some time to get familiar with all the supportive codes.
- **Homework 2 is due on April 19th, Wednesday, 11:59pm, 2017.**